

Report

10570155, 10696253, 10701983

26/04/2021

MATH501 Coursework

Will please add some fancy description here

Machine Learning

Solving tasks in the Machine Learning Section implied possessing basic knowledge of decision trees and KNN as well as the ability to perform Principal Component Analysis. In most cases, the solutions were provided in practical exercises.

Part (a)


Present the data visually using box-and-whisker plots with a distinction for churn. Comment on the data in the context of the problem.

Reading our data in a dataframe:

```
data_path <- "data/churndata.txt"
churn_data <- read.csv(data_path, sep = " ")
churn_data <- na.exclude(churn_data) # removing entries with NA values
```


Plots:

```
churn_data %>% ggplot(aes(x = churn, y = upload, color = churn)) +
  geom_boxplot() +
  labs (y = "Speed of upload",
        x = "Switched to a new provider",
        color = "Switched")
```




Boxplot with average speed of upload against an indicator whether a customer switched to a different provider. As we can see the customers who have not yet switched to a different operator have lower average speed of upload in the internet in comparison with the customers who have switched. In general, having higher uplink speed is a plus since it improves the Internet phone/video call experience and it does no harm to the customers. Hence increasing uplink speed could not affect the customers' decision to switch to different operators.

```
churn_data %>% ggplot(aes(x = churn, y = webget, color = churn)) +
  geom_boxplot() +
  labs (y = "Average time to load a webpage",
        x = "Switched to a new provider",
        color = "Switched")
```



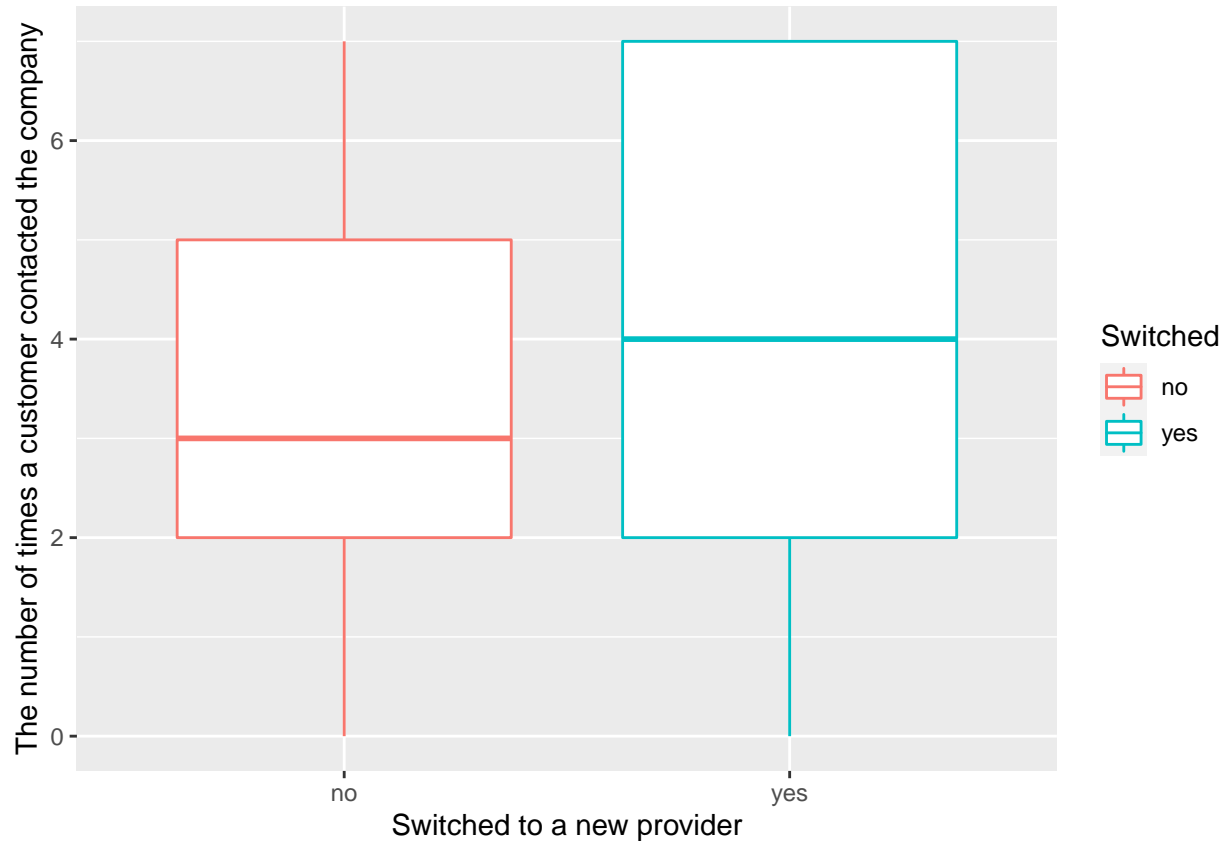
Boxplot with the mean time to load a webpage against an indicator whether a customer switched to a different provider. We can observe a strong dependency between the time to load a webpage (which directly corresponds to the downlink speed) and an indicator whether customers changed their operators. The average downlink speed was significantly lower for the customers that have switched to a different provider than for those who haven't. We can conclude this as the average time to load a webpage for those clients who switched is nearly 200 units longer than of those who didn't.

```
churn_data %>% ggplot(aes(x = churn, y = callwait, color = churn)) +
  geom_boxplot() +
  labs (y = "Customer service waiting time",
        x = "Switched to a new provider",
        color = "Switched")
```



Boxplot with how long a customer waited on the phone call for a customer service operator against an indicator whether a customer switched to a different provider. Even though the average waiting time for a customer service operator is similar in both cases, overall the majority of of customers who switched to a different operator had to wait longer than the average and the customers who haven't changed their provider. We can assume that the time spent by a customer on a call while they're waiting for a customer service operator to attend may impact their decision to switch to another operator although the influence seems to be less significant compared to time to load a webpage.

```
churn_data %>% ggplot(aes(x = churn, y = enqcount, color = churn)) +
  geom_boxplot() +
  labs (y = "The number of times a customer contacted the company",
        x = "Switched to a new provider",
        color = "Switched")
```



Boxplot with the number of times a customer contacted the company via a phone call against an indicator whether a customer switched to a different provider. We can observe that in average the customers that switched to a different operator contacted the company via a phone call 1 time more often than others. The biggest number of contact attempts is 2 calls more than of the customers who haven't changed their providers. Needs to be mentioned that some of the customers who switched didn't contact the company even once. On the contrary, minority of the customers who haven't changed their provider also have more than 5 and even 6 calls. Still, it will be safe to assume that the number of calls impacts the customers' decision to choose a different operator, but its importance is smaller than the time to load a webpage.

Conclusion

Out of all the 4 factors that can possibly influence the 'churn' variable, time to load the webpage (which subsequently leads to the downlink speed) is the most important one. Average phone call customer service waiting time doesn't differ drastically but is still higher for customers who chose different providers; hence, we could conclude that this aspect also plays its part in the customers' decision as well as the number of phone calls to the company. The most suspicious variable is the upload speed - for those clients who changed their providers, the uplink speed was actually higher but the downlink speed was lower (comparing to the customers who didn't change their provider) while normally the opposite should be the case (unless we're talking about 5G). Unfortunately, we don't have access to any other data; hence, we can only speculate that perhaps there are some issues with the provider's network.

Part (b)

Create a training set consisting of 350 randomly chosen data points and a test set consisting of the remaining 150 data points.

```
set.seed(1) # to make the results reproducible
num_subset <- sample(nrow(churn_data), 350) # randomly choose 350 numbers out of 500
```

Separating the data into predictors (X) and classifier (Y):

```
attach(churn_data)

X <- cbind(upload, webget, enqcount, callwait)
Y <- churn
Y <- as.integer(Y == 'yes')
Y <- as.factor(Y)

detach(churn_data)
```

Dividing the data into training and testing sets:

```
train.X <- X[num_subset, ] # 350 records
train.Y <- Y[num_subset]

test.X <- X[-num_subset, ] # 150 records
test.Y <- Y[-num_subset]
```

Part (c)

Using the training data set apply the K nearest neighbours method to construct a classifier to predict churn based on the four available predictors. Find the optimal K using leave-one-out cross-validation for the training data set. Calculate the test error for the classification rule obtained for the optimal K .

Before actually applying KNN we need to normalise the dataset because our data has different units and is on a different scale.

Writing a function for normalising our predictors: subtracting mean value and dividing the result by the standard deviation.

```
normalise <- function (inList){
  m <- mean(inList)
  s <- sd(inList)
  inList <- (inList - m)/s
  return(inList)
}
```

Applying this function to each of the columns in our test and training sets.

```
train.X <- apply(train.X, 2, normalise)
test.X <- apply(test.X, 2, normalise)
```

Now our predictors have values in the same range.

```
head(train.X)
```

```
##      upload      webget  enqcount  callwait
## [1,]  0.9790438  1.5490048 -0.6603278 -0.2912318
## [2,] -1.0678586 -0.7647049  1.7222982 -0.7052967
## [3,] -0.9637788 -1.0506624 -0.6603278  0.5692468
## [4,]  0.2851786 -0.6131415 -0.6603278 -0.1391924
## [5,] -1.4841777 -1.2034098 -0.1838026  0.2878121
## [6,]  0.2851786 -0.3704033  0.7692478  1.4167860
```